An ordered latent field for a cloud-resolving squall line, and what a PDE for its leading channel says

9 November 2026

Abstract

We emulate two-dimensional CM1 squall lines with a *local* latent model: each 2 km patch of the domain becomes a site carrying 32 latent channels, and a stencil law advances the field one minute at a time. Training windows are centred on the storm half of the time, and a structured (nested) dropout orders the channels by importance, so that channel 0 alone reconstructs the full 434-variable column state to nMSE = 0.30, where an ordinary latent needs about 12 channels. Channel 0 is a cold-pool/convection coordinate. Sparse regression on it finds that its evolution is transport, not local growth: a law with spatial derivatives explains 21–42% of its held-out tendency from channel 0 alone, a local ODE explains nothing. Rewritten in physical units through a fixed map to surface θ', the leading law is an advection whose speed is linear in the cold-pool temperature deficit, $c(\theta') = 0.20 + 0.049\,\theta'\,\mathrm{km\,min^{-1}}$, with a tight and test-reproducible slope, and an anti-diffusive term. Its speed at the gust front matches the measured cold-pool speed on average, and a second, independently trained autoencoder gives the same law within about 20%. Three checks of physical meaning fail, however: the speed does not track storm-to-storm differences, its dependence on θ' is 3.5× steeper than density-current scaling, and fitting the same library directly to the surface θ' field gives a slope twice as large, so the coefficient belongs to the latent coordinate, not to θ'. We then ask what the trained propagator does to physical fields, against paired CM1 runs (an exact restart, with and without a small perturbation). It gets the early direction and speed of a perturbation's spread roughly right, but its response is weak (largely the autoencoder) and its leading Lyapunov exponent is 0.006–0.015 $\mathrm{min^{-1}}$ against CM1's 0.12–0.25. Fine-tuning on CM1's perturbation growth raises it on held-out storms but costs forecast accuracy.

1 Data and model

Data. CM1 r21.1 two-dimensional (x–z) squall lines, 400 km wide with $\Delta x = 250$ m, 62 stretched levels, open lateral boundaries and a grid moving at 12.5 $\mathrm{m\,s^{-1}}$; Morrison two-moment microphysics; 6 h runs with 1-min output. 121 storms from three ensembles (control sounding with random initial noise; and boundary-layer moisture × wind-shear variations), split by storm into training, validation and test sets. Each column holds 7 variables (u', w, θ', q_v', p', $\sqrt{q_{\mathrm{liq}}}$, $\sqrt{q_{\mathrm{ice}}}$) on 62 levels, as perturbations from the far-field initial column, standardized per variable: 434 numbers per grid column, 1 600 columns.

Channels and patches. The autoencoder is a 1D convolution along x with the column as the channel dimension (Fig. 1):

- **Patchify**: a convolution with kernel = stride = 8 maps each 8-column patch (2 km) to a hidden vector of width 256. The domain becomes $P = 200$ *sites*.
- **Mix**: two residual convolutions of radius 1 along the sites, so each site sees ±2 neighbouring sites (±4 km); no circular padding (the boundaries are open).
- **Project**: a 1×1 convolution without bias to $c = 32$ latent channels per site.

The latent state is therefore a *field* $z_j(t) \in \mathbb{R}^{32}$, $j = 1..200$, not a single vector. Shifting the input by a whole patch shifts the latent by one site. The decoder mirrors the encoder (transposed convolution, kernel = stride = 8).

Propagator: an explicit latent PDE. The model used here advances the field with a local law, one frame (1 min) per step, integrated with a Heun step:

$$\partial_t z = A_0 z + A_1 \partial_x z + A_2 \partial_x^2 z + b + Q_1[z_i z_j] + Q_2[z_i\, \partial_x z_j] + R(z, \partial_x z, \partial_x^2 z), \tag{1}$$

with centred differences over sites, a small pointwise MLP residual R, and a soft tendency cap. All coefficients start at zero (identity step). The residual ends up carrying only 1–3% of the tendency, so this 32-channel law forecasts CM1 as well as a neural stencil network. But it is dense: tens of thousands of its 52 736 coefficients are needed, and no readable law can be pruned from it. That is the motivation for Section 4: ask for a short law for the *leading* channels of an *ordered* latent instead.

Training. Stage A trains the autoencoder (reconstruction, plus a log-determinant term against channel collapse). Stage B trains the propagator on latent rollouts of up to 45 steps with the autoencoder frozen. Stage C fine-tunes both jointly on decoded 45-step rollouts, with the 3-site halo of every window forced by the true state.

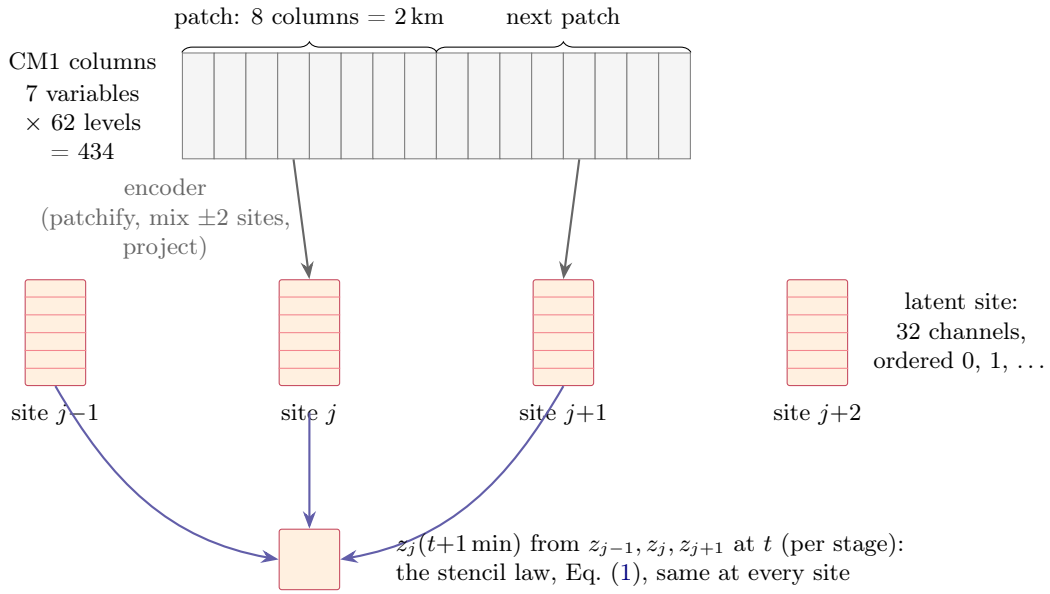


Figure 1: The local latent model. Each 2 km patch of CM1 columns is encoded into one site with 32 channels; nested dropout orders them so that channel 0 carries most of the column. A shared stencil law advances every site from its neighbours.

2 Storm-centred training

Most of the 400 km domain is quiet environment most of the time: a squall line occupies a few tens of kilometres. Uniformly sampled windows would spend most of the training signal on near-zero perturbations. Every stage therefore draws half of its training windows (`active_frac` = 0.5) centred on a storm column, a randomly chosen grid column and time with composite reflectivity above 10 dBZ; the other half is drawn uniformly, so the quiet environment and the storm's surroundings are still

learned (Fig. 2). Windows are 32 sites (64 km) for the autoencoder, 48 sites (96 km) × up to 47 frames for the propagator. Because the model is local, a window is a complete training example: the same weights are applied to the whole 200-site domain at test time.

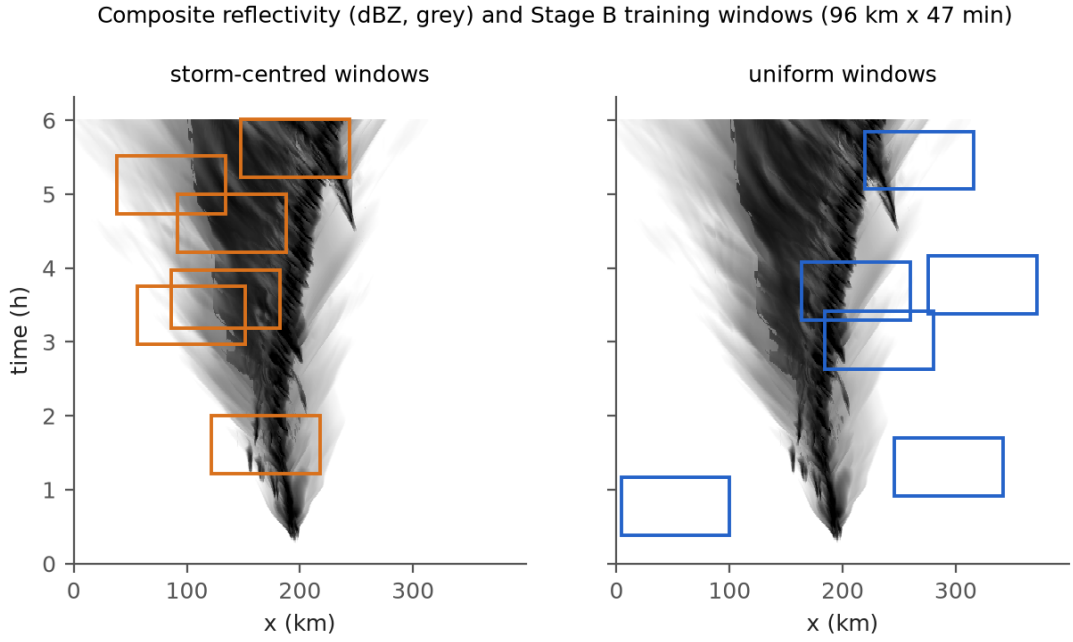


Figure 2: Composite reflectivity of one training storm (grey) with six storm-centred (left) and six uniform (right) Stage B windows. Half of all windows are storm-centred; the uniform half covers the environment.

3 Structured (nested) dropout: an ordered latent

In an ordinary autoencoder the 32 channels are interchangeable: any rotation of the latent is equally good, so no single channel means anything. Nested dropout breaks this symmetry. For half of the training samples, an integer k is drawn uniformly from 1..32, channels k..31 are set to zero, and the decoder must reconstruct the column from channels 0..$k-1$ alone. Channel 0 is then trained to be the best one-channel description of the column, channels 0–1 the best pair, and so on: a nonlinear analogue of a principal-component ordering. The truncation acts on the reconstruction terms of Stage A and Stage C; the propagator always sees all channels.

The ordering is strong (Fig. 3): channel 0 alone reconstructs the column to nMSE = 0.30, and 4 channels to 0.14. The ordinary latent needs about 12 channels to reach 0.5 with its first channels. With all 32 channels the two latents reconstruct about equally well (0.031 against 0.026), so the ordering costs little. A second nested autoencoder trained from scratch with a different seed gives the same curve to within 0.01.

What channel 0 is. On the test storms, channel 0 correlates +0.80 with surface θ', −0.86 with column condensate, −0.72 with column-maximum w and −0.46 with low-level u' (site means). Decoding a field in which only channel 0 is set shows what it encodes (Fig. 4): negative values are a mature convective column, with a surface cold pool of −8 K, latent warming aloft, liquid water below the melting level and an ice anvil above, and a reversal of u' with height (low- and mid-level flows of opposite sign, $\sim$10 m s^{-1}). Positive values are the undisturbed environment. In space and time

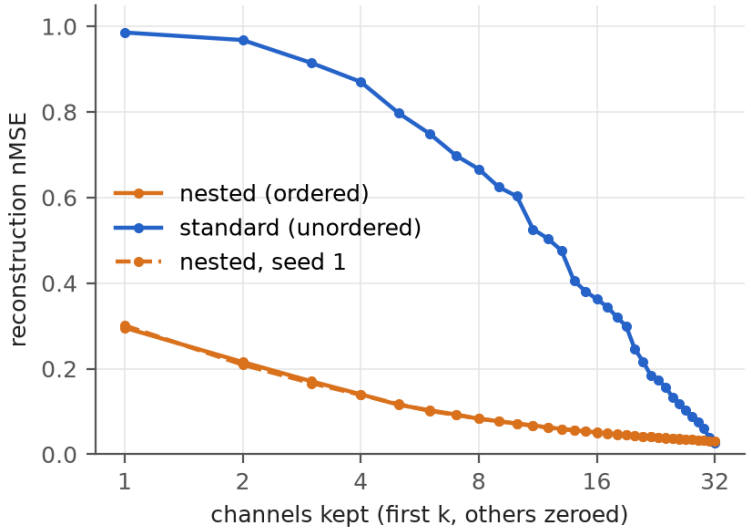


Figure 3: Reconstruction error with only the first k channels kept (others set to zero), on unseen test storms. Nested dropout orders the channels by importance; the ordinary latent spreads information evenly.

(Fig. 6, first two panels) channel 0 tracks the cold pool and the gust front. It is a cold-pool/convection coordinate. Channels 1–3 have no comparably clear physical correlate ($|r| \leq 0.46$).

4 Interpretable PDEs for the leading channels

Method. For the first k channels of the ordered latent ($k = 1, 2, 4, 8$) we fit $\partial_t z_j$ (centred difference, per minute) for $j < k$ by bootstrap sequentially thresholded least squares (STLSQ, resampling whole storms) on 30 training storms, frames 30–330 min, and score on 10 test storms. The candidate libraries are built from the first k channels only:

- **local ODE**: z and quadratic products (no neighbours);
- **PDE**: plus $\partial_x z$, $\partial_x^2 z$ (and optionally up to $\partial_x^4 z$) and their products with z;
- options: products of two derivatives (KPZ-type), cubic terms, second order in time, and a BBM-type mixed operator $(1 - \ell^2 \partial_x^2)\partial_t z$.

Coefficients are in sites (2 km) and minutes. Decision rules for each test were written down before the corresponding run.

Channel 0 is transported, not locally driven. Fig. 5 shows the held-out R^2 of channel 0's tendency. A local ODE explains essentially nothing at any k ($R^2 \leq 0.2$); spatial derivatives explain 21% from channel 0 alone (42% with third and fourth derivatives), rising to 56–61% with two channels and 76–82% with eight. On the ordinary (unordered) latent the same procedure finds almost nothing ($R^2 \leq 0.27$): the ordering is what makes a leading-channel law possible.

The shortest laws. With channel 0 alone, the sparsest PDE (3 terms) is a nonlinear advection,

$$\partial_t z_0 = -(0.059 + 0.082\, z_0)\, \partial_x z_0 - 0.004. \tag{2}$$

With more channels, channel 0's equation keeps the same form, but its advecting velocity is built from several channels: for $k = 8$ its 13 terms are almost all of the form $z_i\, \partial_x z_j$, led by $\partial_t z_0 \approx$

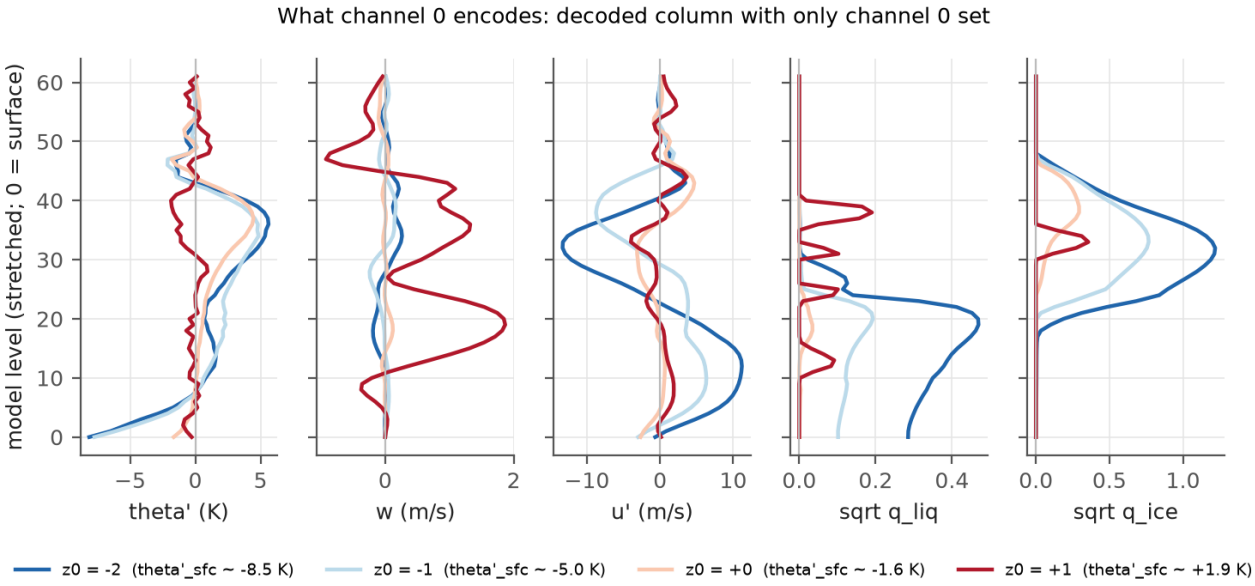


Figure 4: Decoded column at a site when only channel 0 is non-zero, for four values of z_0; the surface θ' each value maps to (Sec. 4.1) is in the legend. Condensate is plotted as $\sqrt{q}$ $(\mathrm{g\,kg^{-1}})^{1/2}$.

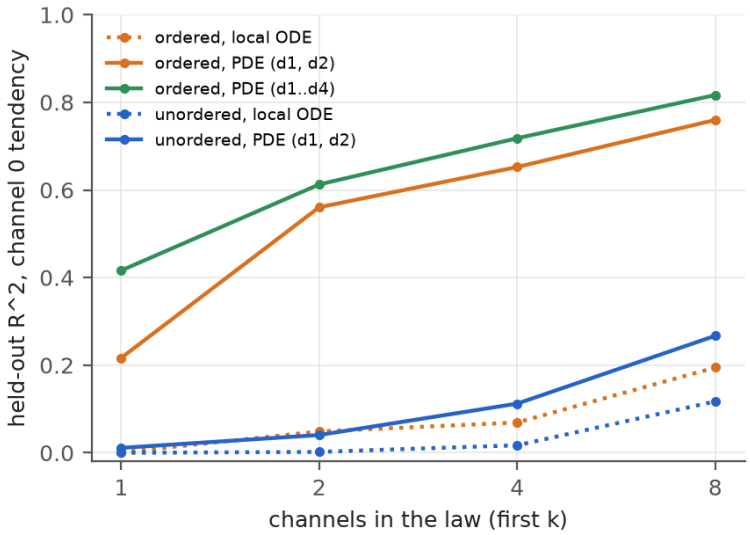


Figure 5: Held-out R^2 of channel 0's tendency against the number of channels in the law (threshold 0.03). Orange/green: ordered latent; blue: ordinary latent. Dotted: local ODE; solid: with spatial derivatives.

$(0.069\,z_1 - 0.064\,z_0)\,\partial_x z_0 - 0.050\,z_1 \partial_x z_1 + \ldots$ (Table 1). The selected structure is stable across bootstrap resamples. Other library extensions add little: a $(\partial_x z_0)^2$ term (KPZ-type) is selected with a consistently negative sign; cubic terms and second order in time do not help; third-derivative terms raise R^2 but their coefficient doubles when the field is smoothed by one more site, so they describe the unresolved sharpness of the gust front rather than a physical dispersion.

Table 1: Selected terms of channel 0's equation (one-step fit, threshold 0.1; coefficients in sites and minutes; bootstrap selection frequency 1.0 unless noted).

law	terms	held-out R^2
$k = 1$, PDE	$-0.082\,z_0 \partial_x z_0$, $-0.059\,\partial_x z_0$ (0.9), -0.004	0.21
$k = 1$, PDE + $(\partial_x z)^2$	$-0.095\,z_0 \partial_x z_0$, $-0.074\,\partial_x z_0$ (0.7), $-0.046\,(\partial_x z_0)^2$ (0.7), … (7 terms)	0.25
$k = 8$, PDE	$+0.069\,z_1 \partial_x z_0$, $-0.064\,z_0 \partial_x z_0$, $-0.051\,z_1 \partial_x z_1$, $+0.026\,z_1 \partial_x z_5$, … (13 terms)	0.65

4.1 The leading law in physical units

A coefficient on z_0 is not physical by itself: the sign, scale and zero point of a latent channel are arbitrary (a gauge). For example, the constant speed in Eq. (2), $0.12\,\mathrm{km\,min^{-1}}$, is the speed at $z_0 = 0$ and changes if z_0 is shifted. We fix the gauge with a physical quantity: on the training storms, $\theta'_{\mathrm{sfc}} = s\,z_0 + b$ with $s = 3.48\,\mathrm{K}$ and $b = -1.55\,\mathrm{K}$ ($r = 0.79$; the OLS and inverse-OLS slopes, 2.76 and 4.38 K, bound the uncertainty of s). Fitting the 4-term law $\partial_t z_0 = k_0 + a_1 \partial_x z_0 + a_2 z_0 \partial_x z_0 + \nu \partial_x^2 z_0$ by ordinary least squares (no sparsity, no tuning for stability) and substituting gives

$$\partial_t \theta' = -c(\theta')\,\partial_x \theta' + \nu\,\partial_x^2 \theta' + S, \qquad c(\theta') = c_0 + c_1 \theta', \tag{3}$$

	R^2	c_0 ($\mathrm{km\,min^{-1}}$)	c_1 ($\mathrm{km\,min^{-1}\,K^{-1}}$)	ν ($\mathrm{m^2\,s^{-1}}$)	c at the front ($\mathrm{km\,min^{-1}}$)
training storms (30)	0.19	+0.202	+0.0490	−614	-0.277 at $\theta' = -9.8\,\mathrm{K}$
bootstrap 95% CI		[0.168, 0.236]	[0.0462, 0.0517]	[−671, −551]	[−0.313, −0.232]
test storms (10)	0.21	+0.173	+0.0476	−628	-0.303 at $\theta' = -10.0\,\mathrm{K}$
seed-1 autoencoder, training	0.24	+0.225	+0.0578	−963	-0.336 at $\theta' = -9.7\,\mathrm{K}$
bootstrap 95% CI		[0.179, 0.273]	[0.0548, 0.0610]	[−1022, −898]	[−0.380, −0.286]
seed-1 autoencoder, test	0.27	+0.190	+0.0567	−940	-0.373 at $\theta' = -9.9\,\mathrm{K}$

The law says (Fig. 7a):

- **The colder the air, the faster it moves in the storm's direction.** The speed is linear in θ' and crosses zero at $\theta' \approx -4\,\mathrm{K}$; at the gust front ($\theta' \approx -10\,\mathrm{K}$) it is -0.28 to $-0.30\,\mathrm{km\,min^{-1}}$, equal to the measured mean speed of the cold-pool centroid (-0.29). The slope c_1 is determined to ±6% and the test storms reproduce it to 3%. A local estimate (the advection speed fitted separately in 2 K bins of θ', grey) is close to the line from -16 to $-6\,\mathrm{K}$ and flattens above.
- **The diffusion is negative** ($\nu \approx -600\,\mathrm{m^2\,s^{-1}}$, CI excluding zero): channel 0's one-step law sharpens the front. As a PDE it is ill-posed, which explains why the closed rollouts of most unrefitted laws blew up, and why refitting the laws on rollouts (not shown) adds large positive diffusion to stay stable. That diffusion is a stability device, not eddy mixing (the physical-field pilot gives $\sim 30\,\mathrm{m^2\,s^{-1}}$).
- **It is one fifth of the story**: $R^2 = 0.2$. The rest of channel 0's tendency needs the other channels (Fig. 5). Fig. 6 compares the law's tendency with the true one: the law captures the advance of the gust front, not the convective cells behind it.

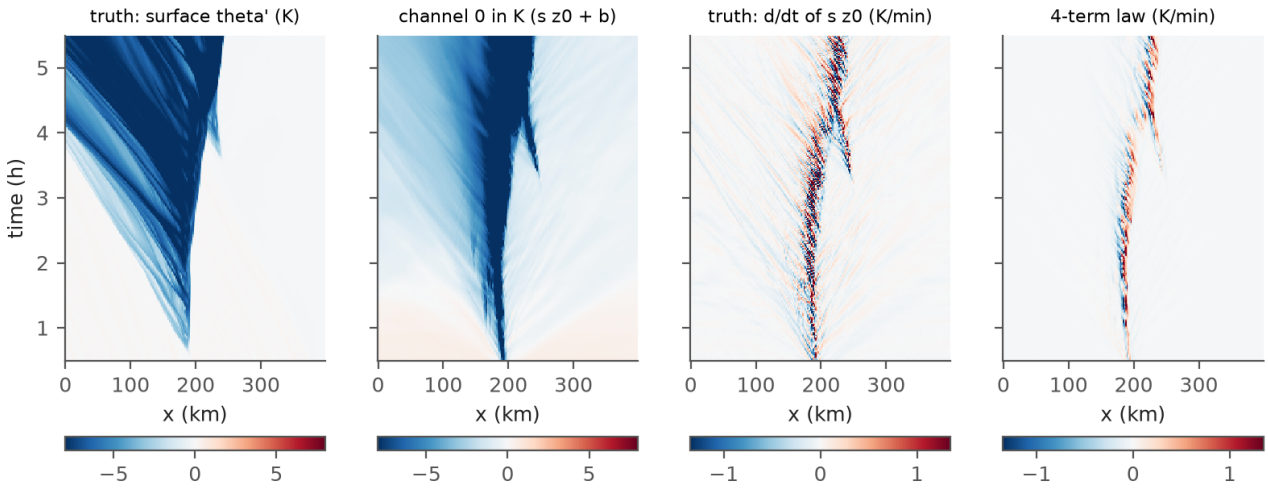


Figure 6: One test storm in x–t. From left: true surface θ'; channel 0 mapped to kelvin; the true tendency of the mapped channel 0; the 4-term law's tendency.

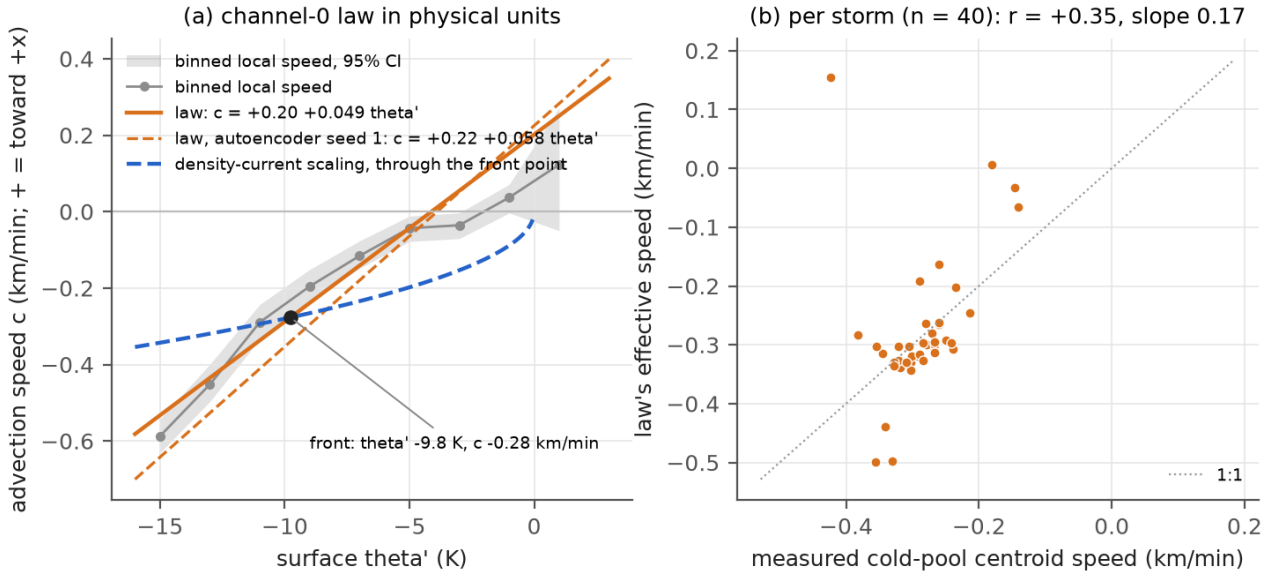


Figure 7: (a) Advection speed of the channel-0 law against surface θ' (orange), the speed fitted locally in θ' bins (grey, bootstrap over storms), and density-current scaling $|c| \propto \sqrt{|\theta'|}$ drawn through the front point (blue, dashed). (b) Per storm: the law's effective speed, fitted on each storm alone, against the measured cold-pool centroid speed.

Does it depend on the autoencoder? A second nested autoencoder, trained from scratch with another seed, is a different latent coordinate system: its channel 0 even has the opposite sign ($\theta'_{\mathrm{sfc}} = -3.46\, z_0 - 1.61\,\mathrm{K}$, $r = -0.81$). After the same gauge fixing, its law has the same form and signs (table, dashed line in Fig. 7a): c_1 is 18% larger and ν 57% more negative, and the two bootstrap intervals for c_1 do not overlap. The law's structure is a property of the data mapped through channel 0; its precise coefficients carry an autoencoder-dependent spread of roughly 20%.

Control: the same library fitted to θ' itself. If the coefficients were properties of CM1, fitting the same four-term library directly to the site-mean surface θ' (same storms, frames, sites and finite-difference target) should give the same numbers. It does not (rule fixed before the run; one autoencoder, one seed):

	R^2 train / test	c_1 (km min^{-1} K^{-1})	ν (m^2 s^{-1})	c at the front (km min^{-1})
θ' directly	0.26 / 0.28	$+0.103$ $[0.096, 0.110]$	-2344	-0.41 $[-0.45, -0.37]$
latent channel 0 (above)	0.19 / 0.21	$+0.049$ $[0.046, 0.052]$	-614	-0.277

The form and signs agree (speed increasing with cold-pool strength, anti-diffusive one-step ν), but c_1 is 2.1× larger on θ' with non-overlapping intervals, and the R^2 is higher. The direct fit overshoots the measured cold-pool speed (-0.29) while the latent's matches it. This does not say which is right: z_0 carries only $r^2 \approx 0.65$ of θ', so errors in variables and the latent's smoothing both bias c_1 low, and I have not separated them. What the control supports is that $c_1 = +0.049$ *is a property of the latent coordinate expressed in kelvin*, and that only the qualitative form is common to both.

Pre-registered tests of physical meaning. Two further tests were fixed before the runs.

1. **Does the law's speed track each storm's speed?** Fitting the advection law on each of 40 storms and comparing with the storm's measured speed: Pearson $r = 0.35$ (cold-pool centroid) and 0.40 (condensate centroid), 95% CIs covering zero, and a slope of 0.2 against the required 0.5–2 (Fig. 7b). **Not met**: the law gets the mean speed right but not the storm-to-storm variation. (With the field smoothed by one site: $r = 0.46$–0.49, slope 0.35; with the seed-1 autoencoder: $r = 0.31$, slope 0.12. Not met.)
2. **Is the θ' dependence density-current-like?** The rule required $c_0 < 0$, $c_1 > 0$ with a CI excluding zero, c_1 within a factor 3 of the density-current prediction $d|c|/d|\theta'| = |c|/(2|\theta'|) = 0.014\,\mathrm{km\,min^{-1}\,K^{-1}}$, the test value inside the training CI, and agreement for an independently trained autoencoder. The sign of c_1, the test agreement and the seed-1 agreement (same signs, c_1 within a factor 2) are met; but c_0 is positive (ambient air moves the other way), and c_1 is 3.5× the density-current value (2.7–4.4× over the bounds on s; 3.3× for seed 1). **Not met.**

4.2 Mixed-derivative (BBM-type) laws

Allowing products of time and space derivatives, $(1 - \ell^2\partial_x^2)\,\partial_t z = F(z)$, with $\ell^2 \in \{0.1, 0.3, 1\}\,\mathrm{sites}^2$ fixed in the sparse fit, two libraries, two thresholds and $k = 1, 2, 4$ (all 36 fits) plus one fit at $k = 8$ (the remaining $k = 8$ fits did not finish within the job's time limit). Three results (Fig. 8).

- **The one-step law is not a good simulator; the rollout refit is.** Closed 30-minute rollouts of the one-step laws beat persistence in only 5 of the 37 fits; after refitting the coefficients to closed rollouts, 35 of 37 do (the exceptions: one fit above persistence, one blow-up), with a z_0 error of 0.067–0.18 against 0.196 for persistence.
- **No systematic gain over the ordinary PDE libraries.** The ordinary libraries reach 0.076–0.19 after the same refit for $k \le 4$. The mixed operator does stabilise one case: the four-derivative library at $k = 4$, threshold 0.03, blows up without it and gives 0.067–0.087 with it.

- **ℓ^2 is a fitting knob, not a length scale.** After the refit, channel 0's ℓ^2 is 0.51 sites^2 at the median (range 0.08–3.6, i.e. $\ell \approx 1.4$ km), is within a factor 2 of its starting value in only 40% of the fits, is uncorrelated with it ($r = 0.2$ on the logarithms) and with rollout skill ($r = 0.09$). (Only channel 0's ℓ^2 was stored; the choice of ℓ^2 is also scored on test storms, so it is exploratory.)

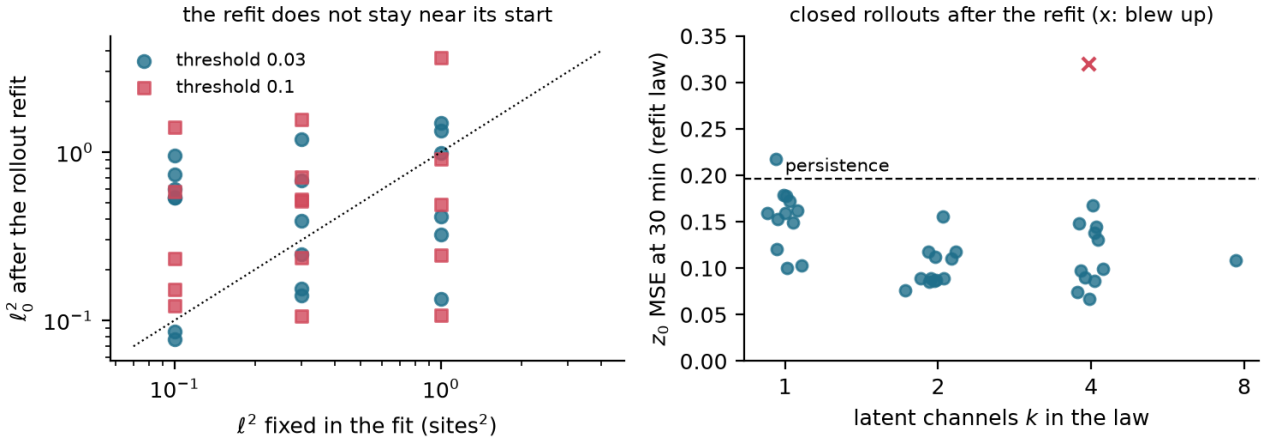


Figure 8: Mixed-derivative laws. (a) Channel 0's ℓ^2 after the rollout refit against the value fixed in the sparse fit; dotted line: unchanged. (b) z_0 error of closed 30-minute rollouts after the refit against the number of channels in the law (dashed: persistence; cross: a blown-up fit).

5 Learning the law together with the propagator

The trained propagator is itself an explicit PDE, but a dense one. Three ways of getting a short law out of it were tried; the decision rules were fixed before each run.

Distilling a law from the propagator. Fitting the four-term channel-0 law to the propagator's own tendency (exact in time, no finite differences) gives held-out $R^2 = 0.19$, the same as against the data (0.21), and 0.26 on states perturbed off the data: the short law's ceiling is not noise but closure. The propagator's channel-0 terms in z_0 alone explain $R^2 = 0.002$ of its channel-0 tendency; channel 0's evolution runs through coupling to other channels. The advection slope is the same as from the data ($c_1 = +0.049\,\mathrm{km\,min^{-1}\,K^{-1}}$, bootstrap 95% interval 0.047–0.052), but the sign of ν flips: $+2900\,\mathrm{m^2\,s^{-1}}$ from the propagator against -614 from the data, so the anti-diffusion of Section 4.1 comes from the finite-difference target.

Forcing channel 0 to be closed. Restricting channel 0's row of the propagator to a short law in z_0 alone, trained jointly with a free encoder, gives a law the data do not follow (data R^2 0.11, against 0.19 before), because rollout training satisfies the constraint by smoothing channel 0 ($\nu \approx 6000\,\mathrm{m^2\,s^{-1}}$) and moving the advection into the other channels; forecast displacement error is 32–37% worse. Not met.

A closed PDE for the four leading channels, with a representation loss. A separate PDE g on channels 0–3 is trained jointly with the dense propagator F by adding $w\,\|F_{<4} - g\|^2/\|F_{<4}\|^2$ to the Stage B and C losses; the autoencoder is free in Stage C. Without the loss, a four-channel library already explains 45% of $F_{<4}$ on held-out storms (channel 0 alone: 19%). With it:

- closure rises to 0.92 ($w = 0.1$) and 0.999 ($w = 1$) (Fig. 9, left);

- but F's agreement with the data's one-step change falls (channel 0: 0.48 to 0.39 and -0.21; channel 2: 0.53 to -0.09 and -0.14), and the 60-minute displacement error rises by 37% and 58% (limit 25%);
- the data-fidelity rule (head against data ≥ 0.4) fails at both weights (mean 0.21 and -0.22).

Closure is therefore bought partly by detaching the model from what the data do: a closed four-channel law for the *model's* leading channels exists, but it is not yet a law for CM1.

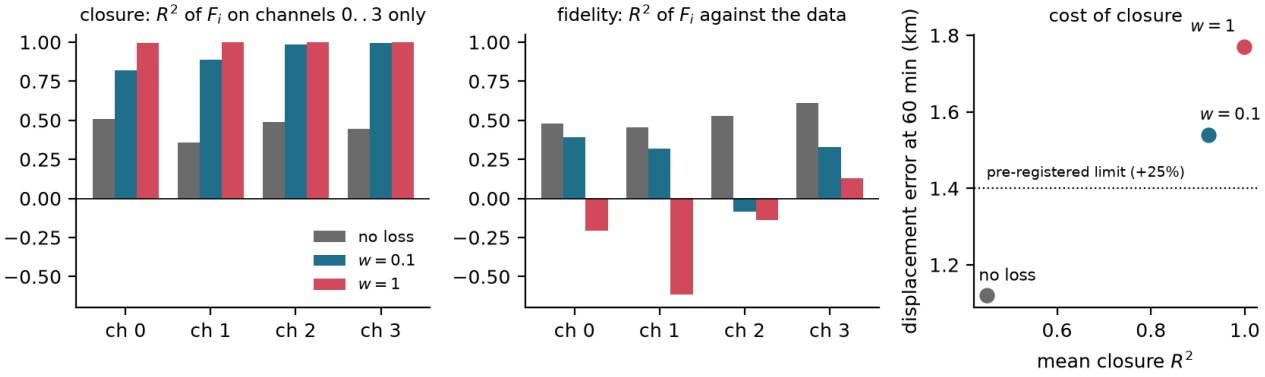


Figure 9: Closure against fidelity for the dense propagator without the loss and with it ($w = 0.1, 1$), test storms. Left: held-out R^2 of each leading channel's tendency on a library of channels 0–3 only. Middle: held-out R^2 of the propagator's tendency against the data's one-step change. Right: 60-minute displacement error against mean closure; dotted: the pre-registered limit.

What the leading channels carry. The correlation of $z_0, \ldots, z_3$ with site means of physical fields (Fig. 10) identifies z_0 as a cold-pool/precipitation-core amplitude (surface θ' +0.82, condensate -0.83, updraft -0.69), z_1 as a surface-outflow pattern (θ' -0.51, u'_{sfc} -0.50), and z_3 as a θ'–q_v pattern; they are partly mixed, so these are labels, not definitions. In the learned four-channel PDE ($w = 0.1$) channel 0's tendency is mostly diffusion and advection of z_0 together with advective coupling to z_1 and z_2 ($z_1\partial_x z_2$, $z_0\partial_x z_2$, $z_0\partial_x z_1$); the z_0-only terms explain only 9.5% of it. The θ' dependence of the speed read from the z_0-only terms is not stable: $c_1 = +0.007$ at $w = 0.1$, -0.005 at $w = 1$.

6 What the propagator does to physical fields

Term-level fitting has reached its limit (Section 5: channel 0's own terms explain 0.2% of the propagator's tendency, so the physics lies in the coupling between channels). A coordinate-free alternative is to ask what decode ∘ propagate ∘ encode does to a physical perturbation, and to check it against CM1.

Truth. For four storms we restart CM1 from a state at 7200 s (the restart reproduces the original run bit for bit) and run 1 h with and without a perturbation: a Gaussian θ' bump (0.5 K, $\sigma = 2$ km, lowest six levels, at $-20, -10, 0, +10$ km from the cold-pool minimum, plus a negative and a 0.05 K one) and a uniform random θ' field of 10^{-3} and 10^{-4} K. CM1's own response is linear at 0.5 K (linearity measure 0.03–0.05, scaling between 0.5 and 0.05 K 0.95–1.00), and its growth from the two tiny amplitudes agrees. The model is given the same bump (encoded, rolled with the control's latents as halo, decoded, minus its own unperturbed rollout).

Results (exploratory; four storms, one seed; two models that differ in latent width):

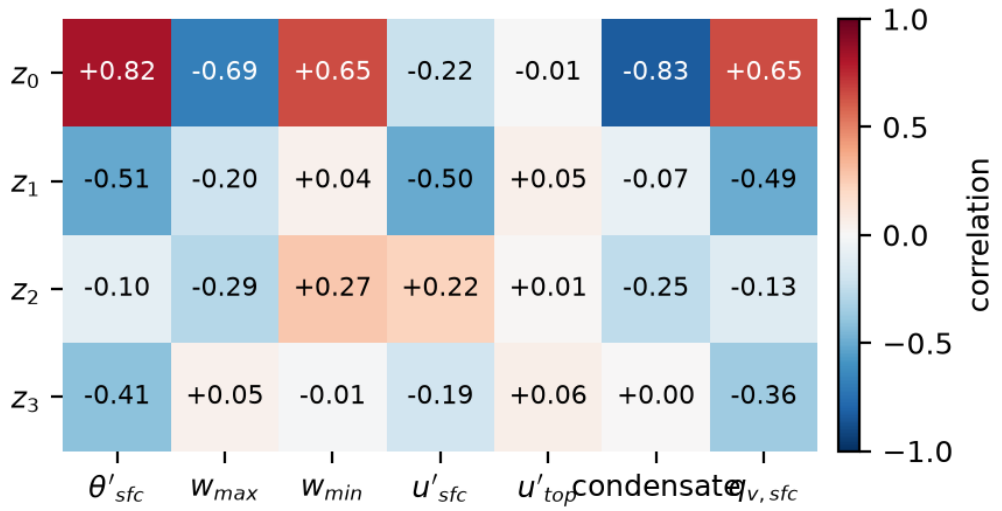


Figure 10: Correlation of each leading channel with site means of physical fields over the training storms (the four-channel model with $w = 0.1$; $w = 1$ agrees to ± 0.04).

	16 channels	32 channels
bump retained by encode/decode at lag 0 (median)	0.047	0.184
response amplitude / CM1 at 15 min	0.053	0.089
response pattern correlation at 15 min	0.52	0.46
leading Lyapunov exponent, 15–30 min (min^{-1}; CM1 0.12–0.25)	0.015	0.006
sensitivity of cold-pool centroid, sign agreement / magnitude ratio	0.92 / 0.05	0.85 / 0.20

- **Direction and early pattern are roughly right.** The response moves toward the storm's direction at 0.6–0.9 $\mathrm{km\,min}^{-1}$ for three of four bump positions, as CM1's does, with pattern correlation about 0.5 at 15 min that decays to near zero by 60 min. The sign of the cold-pool centroid's sensitivity to a bump (ahead of or behind the core) follows CM1 in 85–92% of cases, and the full cold-pool deficit's in 93–100% (a weak test, since any warm bump reduces a cold-pool deficit). Its position of largest sensitivity is not near the cold-pool core and I cannot explain it.
- **The amplitude shortfall is mostly the autoencoder.** The response and sensitivity magnitudes are 5–20% of CM1's. Widening the latent to 32 channels raises the retention 3.9× and the response ratio by 21% less than proportionally; per unit of retained bump the propagator transmits about 0.5–1 of CM1's response in both models.
- **The propagator is an order of magnitude too stable.** Its leading finite-time Lyapunov exponent over 15–30 min is 0.006–0.015 min^{-1} against 0.12–0.25 in CM1, and negative over 0–15 min. This does not change with channel count. A deterministic ensemble of this model should therefore be under-dispersive; I have not tested whether stochastic forcing or inflation repairs this.

Training growth in. We fine-tune only the propagator (autoencoder frozen) so that the growth of the decoded difference between a CM1 branch and its control over 12-minute windows matches CM1's (log ratios, so the autoencoder's amplitude floor drops out), with an anchor term on the decoded control rollout. Training on two storms and testing on the other two (and the swap): the leading exponent over 15–30 min rises on every held-out storm, to 0.06–0.135 min^{-1} on three of four, with the timing of CM1's growth (concentrated in 15–30 min; still negative over 0–15 min). But the pre-registered criterion is not

met: the held-out window loss falls only 44–45% (the training-storm loss to 0.01, memorisation), the held-out control forecast error rises 23–30%, one storm stays at -0.003, and the response amplitude is unchanged. Four storms are too few to say whether more data would fix both; more paired runs are requested.

7 Summary

- A local latent field (32 channels per 2 km site) with storm-centred training emulates CM1 squall lines; its own dynamics is an explicit, but dense, 32-channel PDE.
- Nested dropout gives an importance-ordered latent at no cost in accuracy. Its channel 0 alone carries 70% of the column variance and is a cold-pool/convection coordinate; without the ordering no channel has a physical meaning.
- The leading channel's evolution is transport. In physical units its shortest law is an advection with a speed linear in the cold-pool deficit, a reproducible slope, the right mean front speed, and a negative diffusion that comes from the finite-difference target (fitted to the propagator it is positive).
- The coefficients do not pass the stricter physical tests: they do not track storm-to-storm speed differences, the θ' dependence is not density-current scaling, and the same library fitted directly to θ' gives a slope 2.1× larger. The law is an effective description of how the latent front moves, in a coordinate-dependent unit; only its form (speed increasing with cold-pool strength) is common to the latent and to θ'. It is not a recovered physical equation.
- Fitting the law to the trained propagator instead of the data gives the same advection slope, a positive diffusion, and no better fit: the short law's limit is closure, not noise. A four-channel law is much more closed (0.45), and a representation loss pushes closure to 0.92–0.999, but at the cost of fidelity to the data and of forecast skill; no law that is both closed and faithful has been found.
- Mixed-derivative laws improve closed rollouts only after a rollout refit and add no systematic gain; their length scale ℓ^2 is a fitting knob.
- Against paired CM1 runs, the propagator gets the early direction and pattern of a perturbation's spread roughly right, its response amplitude is mostly limited by the autoencoder, and it is an order of magnitude too stable (leading Lyapunov 0.006–0.015 against 0.12–0.25 min^{-1}). Fine-tuning on CM1's growth raises it on held-out storms but costs 23–30% in forecast error; four storms are too few to settle whether that is fixable.